

Coloured Petri Nets: A Graphical Language for Formal Modeling and Validation of Concurrent Systems

Kurt Jensen
Department of Computer Science
Aarhus University, Denmark
kjensen@cs.au.dk

Lars M. Kristensen
Department of Computing
Bergen University College, Norway
lmkr@hib.no

ABSTRACT

Coloured Petri Nets (CPNs) combine Petri nets with a programming language to obtain a scalable formal modeling language for concurrent systems. Petri nets provide the formal foundation for modeling concurrency and synchronization, and a programming language provides the primitives for modeling data manipulation and creating compact and parameterizable models. We provide an example driven introduction to the core syntactical and semantical constructs of the CPN modeling language, and briefly surveys how quantitative and qualitative behavioral properties of CPN models can be validated using simulation-based performance analysis and explicit state space exploration. In addition, we give a brief overview of CPN Tools which provide tool support for the practical use of CPNs, and provide pointers to some significant examples where the CPN technology has been put into practical use in an industrial setting. As we proceed, we provide a historical perspective on the research that led to the development of the CPN language.

Categories and Subject Descriptors

H.4 [Information Systems Applications]: Miscellaneous;
D.2.8 [Software Engineering]: Metrics—*complexity measures, performance measures*

General Terms

Theory

Keywords

TODO

1. INTRODUCTION

The vast majority of IT systems today can be characterized as concurrent and distributed systems in that their operation inherently relies on communication, synchronization, and resource sharing between concurrently executing software components and applications. This is a development that has been accelerated first with the pervasive presence

of the Internet as a communication infrastructure, and in recent years by, e.g., cloud- and web-based services, mobile applications, and multi-core computing architectures.

The development of Coloured Petri Nets (CPNs) was initiated in the early 80's when distributed system was considered to become a major paradigm for future computing systems. The goal of the CPN modeling language was to develop a formally founded modeling language for concurrent system that would make it possible to formally analyze and validate concurrent systems, and which from a modeling perspective would scale to industrial systems. A main motivation behind the research into CPNs (and many other formal modeling languages) was that the engineering of correct concurrent systems is a challenging task due to their complex behavior which may result in subtle bugs. As concurrent systems are becoming still more pervasive and critical to society, formal techniques for concurrent system was – and still is – a highly relevant technology to support the engineering of reliable concurrent systems.

At its very base, CPNs builds on the visionary work of C. A. Petri [?] who already in the 60's introduced Petri Nets as a formalism for concurrency and synchronization. In Petri Nets, concurrency is a fundamental concept in that Petri Nets is inherently based on the idea that everything is (implicitly) concurrent unless explicitly synchronized. This is in contrast to many other modeling formalisms where concurrency must be explicitly introduced using parallel composition operators. A further advantage of Petri nets is that they rely on very few basic concepts, and is still able to model a wide range of communication and synchronization concepts and patterns. A disadvantage of Petri nets in their basic form is, however, that they do not scale to large systems unless one models the systems at a very high level of abstraction. The primary reasons for this is that Petri nets are not suited for modeling sequential computation and data manipulation and they do not provide concepts that make it easy to scale models according to some system parameter, e.g., increase the number of servers in a modeled system without having to make major changes to the model.

2. COLOURED PETRI NETS

To present the concepts of CPNs, we use a CPN model of a distributed two-phase commit transaction system. We use the example to give an informal introduction to CPNs. The reader interested in the formal definition of CPNs is referred to [?].

The CPN model of the two-phase commit system is comprised of 4 *modules* hierarchically organized into three levels. Figure 1 shows the top-level module which consists of two *substitution transitions* (drawn as rectangles with double-lined borders) representing the **Coordinator** and the **Workers** in the system. Each of the substitution transitions has an associated *submodule* that model the behavior of the coordinator and the workers, respectively. The name of the submodule is written in the rectangular tag positioned at the bottom of each substitution transition.

The two substitution transitions are connected via directed *arcs* to the four places **CanCommit**, **Votes**, **Decision**, and **Acknowledge** (drawn as ellipses). Places connected to substitution transitions are called *socket places* and are linked to *port* places on the associated submodules (to be presented shortly). The coordinator and the workers interact by producing and consuming *tokens* on the places. These tokens carry data values and the type of tokens that may reside on a place is determined by the *type* of the place (written in text below the place). For historical reasons the types of places are called *color sets*. Figure 2 lists the definitions of the color sets used for the four places in Fig. 1. These color sets are defined using the CPN ML programming language.

The **Worker** color set is an indexed type consisting of the values `wrk(1)`, `wrk(2)`, ..., `wrk(W)` where the symbolic constant `W` is used to specify the number of worker processes considered. This color set is used to model the identity of the worker processes. The color set **Vote** is an enumeration type containing the values **Yes** and **No**, and is used to model that a worker may vote **Yes** or **No** to commit the transaction. The color set **WorkerxVote** is a product type containing pairs consisting of a worker and its vote. The color set **Decision** is an enumeration type used to model whether the coordinator decides to **abort** or **commit** the transaction (only if all workers vote yes will the transaction be committed). It should be noted that in addition to the type constructors introduced above, CPN ML supports union, lists, and record types. The variables `w` and `vote` declared in Fig. 2 will be introduced later.

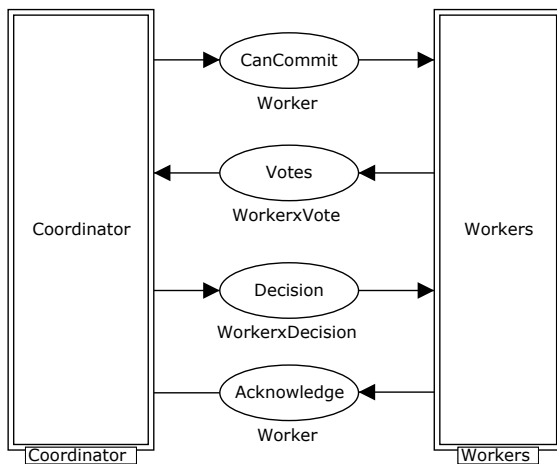


Figure 1: The top-level module of the CPN model.

```
val W = 2;

colset Worker = index wrk with 1..W;
colset Vote = with Yes | No;
colset WorkerxVote = product Worker * Vote;

colset Decision = with abort | commit;
colset WorkerxDecision = product Worker * Decision;

var w : Worker;
var vote : Vote;
```

Figure 2: Definition of color sets used in Fig. 1.

The state of a CPN model is called a *marking*, and consists of a distribution of *tokens* on the places of the model. Each place may hold a (possibly empty) *multi-set* of tokens with data values (colors) from the color set of the place. The *initial marking* of a place is specified above each place (and omitted if the initial marking is the empty multi-set). For the places in Fig. 1 all places are empty in the initial marking. Initially, the *current marking* of a CPN model equals the initial marking. When a CPN model is executed occurrences of enabled transitions consume and produce tokens on the places which will change the *current marking* of the CPN model.

The **Coordinator** module is shown in Fig. 3. This is the submodule associated with the **Coordinator** substitution in Fig. 1. The places **CanCommit**, **Votes**, **Decision**, and **Acknowledge** are *port places* as indicated by the rectangular **In** and **Out** tags positioned next to them. These places are linked to the accordingly named places in the top-level module (Fig. 1) via a *port-socket association* which implies that any tokens added (removed) from a port place by transitions in the **Coordinator** module will also be added (removed) in the marking of the associated socket place in the top-level module. The places **Votes** and **Acknowledge** are *input port places* which means that the module will only consume tokens from these places. The places **CanCommit** and **Decision** are *output port places* which means that the **Coordinator** module will only produce tokens on these places. It is also possible for a place to be an input-output port place which means that the module may both consume and produce tokens from (on) this place.

The places **Idle**, **WaitingVotes**, and **WaitingAcknowledgement** are used to model the states that the coordinator goes through when executing the two-phase commit protocol. The places **Idle** and **WaitingVotes** have the color set **UNIT** containing just a single value `()` (denoted unit). Initially, the coordinator is in an idle state as modeled by the initial marking of place **Idle** which consists of a single token with the color unit. In CPN ML this multi-set is written `1'()` specifying one (1) occurrence of `()` the unit color `()`. The number of tokens on a place in the current marking is indicated with a small circle positioned next to a place, and the detail of the color of the tokens are provided in an associated text box. The indication of the current marking of a place is omitted if currently the place contains no tokens.

The transitions **SendCanCommit**, **CollectVotes**, and **ReceiveAc-**

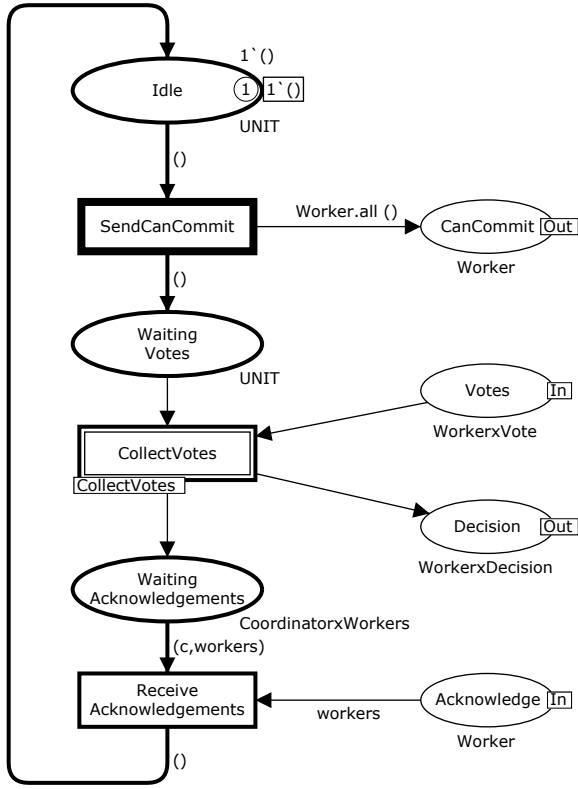


Figure 3: The Coordinator module.

knowledge model the events/actions that cause the coordinator to change state. The coordinator will first send a can commit message (transition **SendCanCommit**) to each worker asking whether they can commit the transaction. Then the coordinator will collect the votes from all workers (transition **CollectVotes**) and send a decision to the workers that voted yes indicating whether the transaction is to be committed or not. Finally, the coordinator will receive an acknowledgment from each worker that voted yes, confirming that they have received the decision (transition **ReceiveAcknowledgements**). It should be noted that **CollectVotes** is a substitution transition which means that the details of how the coordinator collects votes is modeled by the associated **CollectVotes** submodule. This illustrates that it is possible to mix the use of ordinary and substitution transitions within a module. We do not describe the **CollectVotes** module in this paper. The complete CPN model is available via [?].

In the current marking shown in Fig. 3 only the transition **SendCanCommit** is enabled as indicated by the thick border of that transition. The requirement for a transition to be *enabled* is determined from the *arc expressions* associated with the incoming arcs of the transition. In this case, there is only a single incoming arc from place **Idle** containing the expression $()$. This expression specifies that for **SendCanCommit** to be enabled, there must be at least one $()$ -token present on **Idle**. When the **SendCanCommit** transition occurs, it will consume a $()$ -token from place **Idle** and it will produce tokens on places connected to output arcs as determined by evaluating the arc expressions on output arcs. In

this case, the expression $()$ on the arc to **WaitingVotes** evaluates to a single $()$ -token. The expression **Worker.all()** is a call to the function **Worker.all()** that takes a unit value $()$ as parameter and returns all colors of the color set **Worker**. This illustrates that it is very useful to be able to hide complex calculations (e.g., of multi-sets of tokens or tokens with complex data values) within function calls and to be able to add/remove several tokens in a single step (transition occurrence) without introducing a number of intermediate states (markings).

Figure 4 shows the marking of the surrounding places of transition **SendCanCommit** after the occurrence of **SendCanCommit**. It can be seen that the place **CanCommit** contains two tokens - one of each worker in the system - representing messages going to the two worker processes. The coordinator has now entered a state in which it is waiting to collect the votes from the worker processes.




Figure 4: Current marking after SendCanCommit.

By setting the symbolic constant **W** (see Fig. 2) we can easily configure the model to handle, e.g., five workers. With ordinary Petri nets we would have had to create a copy of the **CanCommit** place for each worker. In particular, we would have to make changes to the net structure (places, transitions, arcs) when changing the number of workers. This shows that CPNs provides a means for easily creating parameterizable models and also that it enables more compact modeling as we only need a single instance of the **CanCommit** place in order to accommodate any finite number of workers.

- WE COULD ELABORATE MORE ON THIS BY DRAWING THE CORRESPONDING PT-NET FOR THE CANCOMMIT TRANSITION. THAT WOULD ILLUSTRATE UNFOLDING OF PLACES (see comment at the end of this section that proposes illustration of unfolding transition. Perhaps it is a good approach to split the introduction of unfolding in two steps. Then at the end of the section we can mention that any CPNs can be unfolded to a possibly infinite behaviorally equivalent PT-nets. The two example that we have in this section would then nicely illustrate how. It would also be possible to factor this into a separate section in order to avoid that this section becomes overlong in comparison with the other section.

Figure 5 shows the **Worker** module which is the submodule

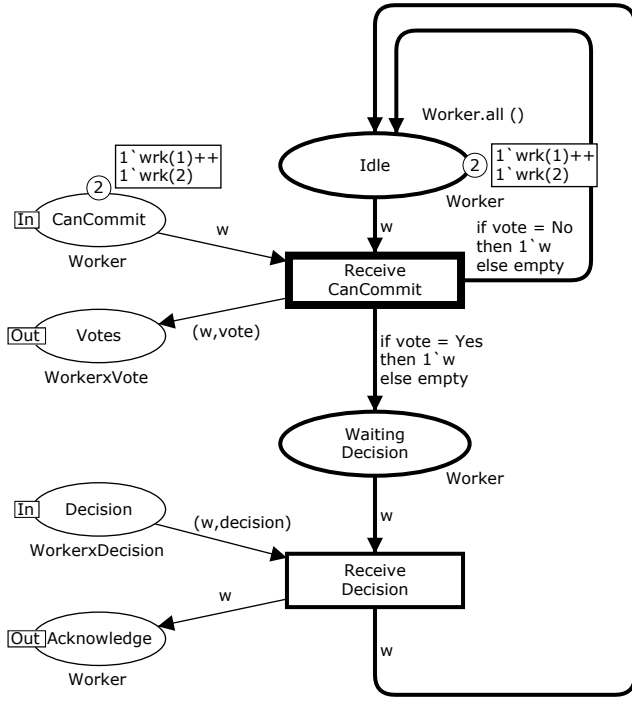


Figure 5: The Worker module.

of the **Workers** substitution transition in Fig. 1. The places **CanCommit**, **Votes**, **Decision**, and **Acknowledge** constitute the port places of this module and are linked to the accordingly named socket places in Fig. 1. The places **Idle** and **WaitingDecision** models the two main states of worker processes. Each of these places have the color set **Worker** and the idea is that when there is a token with color **wrk(i)** on, e.g., the place **Idle**, then this models that the *i*'th worker is in state idle. This makes it possible to model the state of all workers in a compact manner within a single module without having to have a place for each worker or a module instance for each worker. Initially, all workers are in the idle state as represented by corresponding tokens on place **Idle** in the initial marking. The transition **ReceiveCanCommit** models the reception of can commit messages from the coordinator and the sending of a vote. The transition **ReceiveDecision** models the reception of a decision message from the coordinator and the sending of an acknowledgment.

The current marking of place **CanCommit** is $1'wrk(1) ++ 1'wrk(2)$ modelling a marking where the coordinator has sent a can commit message to each worker. The thick border of transition **ReceiveCanCommit** indicates that this transition is enabled in the current marking. The arc expressions on the surrounding arcs of the **ReceiveCanCommit** transition are more complex than the arc expressions of the **SendCanCommit** transition in the **Coordinator** module considered earlier in that they contain the **free variables** **w** and **vote** defined in Fig. 2. This means that in order to talk about the enabling and occurrence of transition **ReceiveCanCommit**, we need to assign values to these variables in order to evaluate the input and output arc expressions. This is done by creating a *binding* which associates a value to each of the free variables occurring in the arc expression of the transition. Bindings

can be considered different modes in which a transition may occur. As **w** is of type **Worker** and **vote** is of type **Decision**, this gives the following four possible bindings reflecting that each of the two workers may vote yes or no to committing the transactions:

$$\begin{aligned} b_{1Y} &= \langle w = wrk(1), vote = Yes \rangle \\ b_{1N} &= \langle w = wrk(1), vote = No \rangle \\ b_{2Y} &= \langle w = wrk(2), vote = Yes \rangle \\ b_{2N} &= \langle w = wrk(2), vote = No \rangle \end{aligned}$$

A binding of a transition is enabled if evaluating each input arc expression in the binding results in a multi-set of tokens which is a subset of the multi-set of tokens present on the corresponding input place. For an example, consider the binding b_{1Y} . Evaluating the input arc expression **w** on the input arc from **Idle** results in the multi-set containing a single token with the color **wrk(1)** which is contained in the multi-set of tokens present on place **Idle** in the marking depicted in Fig. 5. Similarly for the input arc expression on the arc from place **CanCommit**. This means that binding b_{1Y} is enabled and may occur. In fact, all four bindings listed above is enabled in the marking shown in Fig. 1.

The tokens produced on output places when a transition occurs in an enabled binding is determined by evaluating the output arc expressions of the transition in the given binding. Consider again the binding element b_{1Y} . The output arc expression $(w, vote)$ will evaluate to $(wrk(1), Yes)$ and this token will be added to place **Votes** to inform the coordinator that worker 1 votes yes to committing the transaction. The arc expression on the arc from **ReceiveCanCommit** to **WaitingDecision** is an if-then-else expression which in the binding b_{1Y} will evaluate to the multi-set $1'wrk(1)$ which will be added to the tokens on place **WaitingDecision**. The if-then-else expression on the arc from **ReceiveCanCommit** to **Idle** evaluates to the **empty** multi-set and hence no tokens will be added to place **Idle** in this case. Figure 6 shows the marking resulting from an occurrence of the b_{1Y} binding.

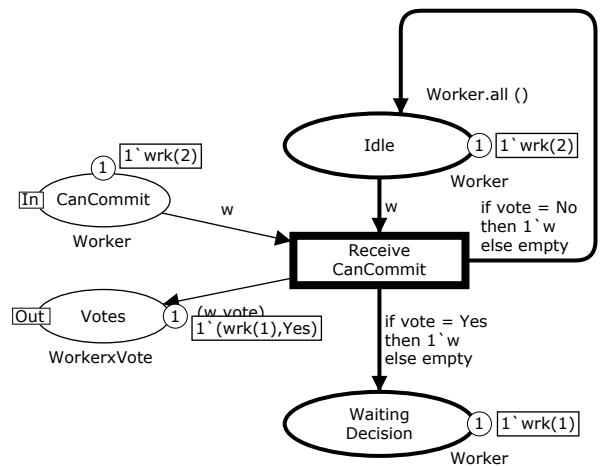


Figure 6: Current marking after occurrence of **ReceiveCanCommit**.

The occurrence of the binding b_{1N} representing that worker one votes no would have the effect of removing a $\text{wrk}(1)$ -token from `Idle`, adding a $(\text{wrk}(1), \text{No})$ -token to `Votes`, and adding no tokens to place `WaitingDecision` and adding a $\text{wrk}(1)$ -token to place `Idle`. This models the fact that if a worker votes no to committing the transaction, then it goes back to idle; whereas if it votes yes, then it will go to waiting to be informed about whether the transaction is to be committed or not. Recall that this is a distributed system and hence a worker cannot (without exchanging messages) know what other workers have voted. Above we have considered relatively simple arc expressions but the arc expressions of a transition can be any expression that can be written in Standard ML as long as they have types that matches the corresponding places. In particular, arc expressions may apply functions including higher-order functions.

All four bindings listed for transition `ReceiveCanCommit` were enabled in the marking shown in Fig. 5. Furthermore, enabled bindings may be *concurrently enabled* if each binding can get its required multi-set of tokens from each input place independently of the other enabled bindings in the set. As an example, the two bindings b_{1Y} and b_{2Y} are concurrent enabled since each binding can get its token from the input places without sharing with each other. This reflects that the workers are executing concurrently and may simultaneously send a vote back to the coordinator. In contrast, the two bindings b_{1Y} and b_{1N} are not concurrently enabled. These two bindings are in *conflict* because they each need the single $\text{wrk}(1)$ -token on `Idle` (and in fact also the single $\text{wrk}(1)$ -token on place `CanCommit`). The notion of concurrency and conflict of binding extends to bindings of different transitions. A fundamental property of a set of concurrently enabled bindings that CPNs inherit from Petri nets is that theses can be executed in any interleaved order and the resulting marking will be the same independently of the interleaved execution considered.

- HERE WE COULD MAKE A FURTHER LINK TO PT-NETS BY SHOWING THE FRAGMENT CORRESPONDING TO RECEIVECANCOMMIT. THAT WOULD ILLUSTRATE UNFOLDING OF TRANSITIONS.
- HERE WE COULD BRIEFLY MENTION GUARDS.
- WE COULD BRIEFLY TALK ABOUT MODELING TIME IF WE SKIP HAVING A SECTION ON THIS.

3. HISTORY OF CPNS

The basic ideas of Petri nets were introduced by Carl Adam Petri in his doctoral thesis published in 1962 [?]. This was far ahead of the time where distributed systems were invented and computers started to have parallel processes. At that time programs and processing were considered to be sequential and deterministic. Hence, it was extremely visionary of Carl Adam Petri to predict the importance of being able to understand and characterize the basic concepts of concurrency.

Over the next decade, Petri nets were widely accepted as one of the most well-founded theories to describe important

behavioral concepts such as concurrency, conflict, synchronization and resource sharing. Petri nets were also used to model and analyze small practical systems. However, the practical use soon revealed a serious shortcoming. Petri nets (in their basic form) are not well-suited to model systems in which data play a crucial part. Since this is the case for most computing systems (and many other kinds of systems) the use of Petri nets for practical modeling were staggering. To remedy this situation many modelers proposed different ad-hoc extensions to Petri nets. This created a large "zoo" of different Petri net modeling languages. Many ad-hoc extensions were not well-defined, and even when they were, they introduced a fundamental problem. Whenever a new ad-hoc extension was introduced, all the basic concepts and analysis methods had to be redefined - to apply for the extended Petri net language (with the ad-hoc extension).

With the invention of the first (text-based) computer tools to support the analysis of Petri net models, the situation became acute. Whenever a new ad-hoc extension was introduced (to handle a modeling problem) all existing computer tools became void, and could only be used after time-consuming (and error-prone) reprogramming/extension. Hence, there was an urgent need to develop a class of Petri nets that were general enough to handle a large variety of different application areas without the need of making ad-hoc extensions.

The first successful step towards a common more powerful class of Petri nets were taken by Hartmann Genrich and Kurt Lautenbach in 1979 with the introduction of Predicate/Transition Nets (PrT nets) [?]. Their work was inspired by earlier work on transition nets with "colored tokens" by M. Schiffers and H. Wedde and transition nets with "complex conditions" by Robert Shapiro. The basic idea behind PrT nets was to introduce a set of colored tokens which can be distinguished from each other - in contrast to the indistinguishable black tokens in basic Petri nets. In this way it became possible to model different processes in a single subnet. As an example, consider the workers in the CPN model of Sect. 2. In the basic Petri net model, it would be necessary to have a subnet for each worker (even though they behave in exactly the same way), but in PrT nets they can be modeled by a common subnet, because we can use tokens with color $\text{wrk}(1)$ to model the state of the first worker, tokens with the color $\text{wrk}(2)$ to model the state of the second worker, and so on. This means that we for W workers can have a single `Idle` place, which may contain tokens of W different colors - instead of having a separate `Idle` place for each of the W workers. PrT nets use arc expressions to define how transitions can occur in different ways (occurrence modes) depending of the colors of the involved input and output tokens. This works in a similar way as described for CPNs in Sect. 2.

The invention of colored distinguishable tokens in PrT nets was a gigantic step forward - but it still had some limitations. PrT nets only had one set of token colors, and all places have to use this set (or Cartesian products of this set with itself). For our example in Sect. 2 this means that the identity workers, Yes/no votes and abort/commit decisions all have to be modeled by colors in this single unstructured set of token colors.

The second step towards a more general class of Petri nets was taken by Kurt Jensen in his PhD thesis 1980 with the introduction of the first kind of Colored Petri Nets [?]. This net model allowed the modeler to use a number of different color sets (e.g. one color set for the coordinator, a second for the workers, a third for Yes/no votes and a fourth for abort/commit decisions). This made it possible to represent data values in a more intuitive way instead of having to encode all data into a single shared set. It later turned out to be convenient to define the color sets by means of data types known from programming languages, such as products, records, lists, and enumerations. The use of types had three implications: Token colors became structured (and hence much more powerful). Type checking became possible (making it much easier to locate modeling errors). Color sets, arc expressions and guards could be specified by the well-known and powerful syntax and semantics known from programming languages. This gave the modeler a convenient way to handle complex data and specify the often complex interaction between data and system behavior.

A third step forward was taken by Peter Hubert, Kurt Jensen and Robert Shapiro in 1990 with the introduction of Hierarchical CPNs [?]. Their work was heavily inspired by the hierarchy concepts in the Structured Analysis and Design-Technique (SADT) developed by D.A. Marca and C.L. McGowan. It was Robert Shapiro who got the idea to port the SADT hierarchy concepts to CPN. Hierarchical CPNs allows the modeler to split a large model into a number of interacting and re-usable modules (components) - in a similar way as used in many programming languages. The basic ideas are illustrated by the modules, substitution transitions and port/socket places in Sect. 2. The introduction of hierarchies in CPNs has several implications: Petri net models of large systems become much more tractable, since they can be split into modules of a reasonable size (instead of working with a single monolithic net which has to be glued to a large wall (or laid out on a football field)). The human modeler can concentrate on a few details (in a single module) at a time - instead of being overwhelmed with the full details of the complete model. CPN modules can be seen as black boxes, where modelers, when they desire, can forget about the details within the modules. This makes it possible to work at different abstraction levels - and hence we talk about hierarchical CPNs. Finally there are often system components that are used repeatedly. It would be inefficient to model these components several times. Instead a module can be defined once and used repeatedly. In this way there is only one description to read, and one description to modify, when changes are necessary.

The fourth step was taken by the creation of graphical computer tools to support the modeling and analysis by means of CPNs. The Design/CPN tool was created at Meta Software, Cambridge, Massachusetts, USA starting in 1988 [?]. The main architects behind the tool were Kurt Jensen, Robert Shapiro and Peter Hubert and the implementation was made together with an international group of people including Jawahar Malhotra (who got the brilliant idea to use the Standard ML language for type definitions and net inscriptions), Ole Bach Andersen (who implemented the graphical interface), S  ren Christensen (who implemented the complex algorithms for automatic binding of free variables dur-

ing simulations) and Hartmann Genrich (who contributed with knowledge and experience from PrT nets). The first version of Design/CPN supported modeling, syntax check and interactive simulation. Later versions added timed CPNs (based on ideas by Wil van der Aalst), fast automatic simulation (e.g. for performance analysis) and state space analysis (to investigate behavioral properties).

Steps 2-4 above are closely related. It was during the creation of the Design/CPN tool that the need of modules was discovered and it was also in this phase that it became clear that the type concept from programming languages was extremely adequate for type definitions and net inscriptions (instead of using more ad-hoc notations). If we compare the development of CPNs with the development of programming languages, we find many similarities. The first step from black tokens to colored tokens corresponds to the step from bits to simple data types. The second step corresponds to the invention of structured data types and type checking. The third step corresponds to the introduction of structuring concepts like modules, procedures, functions and subroutines. The fourth step corresponds to the implementation of compilers and run-time systems. Without these programming languages are of no practical use. Analogously, a Petri net language without tool support is useless for practical modeling work.

The extensions of Petri nets and programming languages described above add modeling power. They make it easier for the user to model/program complex systems. The extensions do not add computational power. Anything which can be programmed in Java can (in theory) also be programmed in assembler code. It is just much more time-consuming and much more error-prone. Analogously it can be proved (quite easily) that each hierarchical CPN can be unfolded to a (much larger) basic Petri net (Place Transition Net) with exactly the same dynamic behavior. Furthermore, each Place Transition Net can be folded into a CPN with a single module consisting of a single place and a single transition. In practice such a folding is totally uninteresting, because the arc expressions will be extremely complex and totally non-interpretable for a human being. However, the fact that the unfolding and folding exists shows that hierarchical CPNs has the same theoretical properties as basic Petri nets - in particular that CPNs are a solid model for concurrency, conflict synchronization and resource sharing.

The introduction of hierarchical CPNs supported by the Design/CPN tool made a dramatic change to the practical use of Petri nets. The new modeling language and its tool support were general and powerful enough to eliminate the need of making ad-hoc extensions. A common platform for practical modeling had been established and this was used by most Petri net practitioners. The use of the platform was supported by a three volume monograph on Colored Petri Nets published by Kurt Jensen in 1992-1997 [?]. In addition a large number of research papers were published. Some of these describe new analysis methods, while others describe experiences from practical modeling and analysis. References to more than hundred modeling/analysis paper can be found in [?]. Many of the projects have been carried out in an industrial environment.

Starting from year 2000 a second generation of tool support, called CPN Tools, was designed and implemented at Aarhus University, Denmark. The main architects behind the new tool were Kurt Jensen, S  ren Christensen and Michael Westergaard. The graphical user interface was designed together with Michel Beaudouin-Lafon and Wendy McKay from the international HCI research community. It is based on empirical studies of the use of Design/CPN and much easier and efficient to use. There is also a much faster simulation engine (developed by Torben Haag and Tommy Hansen). Many models run one thousand times faster allowing complex automatic simulations to be executed within seconds instead of hours. Many different kinds of state space analysis are supported (developed by Lars Kristensen, Michael Westergaard and Thomas Mailund). There is high-level support for defining and collecting data from simulation-based performance analysis (developed by Lisa Wells and Bo Lindstr  m). Finally there is improved support for creating graphical feedback from ongoing simulations (developed by Michael Westergaard).

An international standard for high-level Petri nets was developed and approved in 2004 [?]. The standardization work was headed by Jonathan Billington and the standard is heavily based on CPNs (which adhere to the standard).

A new monograph on Colored Petri Nets has been published by Kurt Jensen and Lars M. Kristensen in 2009 [?]. It provides an in-depth, but yet compact introduction to modeling and validation of concurrent systems by means of CPNs. The book introduces the constructs of the CPN modeling language, presents its analysis methods, and provides a comprehensive road map to the practical use of CPNs. Furthermore, the book presents some selected industrial case studies illustrating the practical use of CPN modeling and validation for design, specification, simulation, and verification in a variety of application domains. The book is aimed at use both in university courses and for self-study.

In 2010 CPN Tools had 10.000 licenses in 150 countries. At that time the development and maintenance of the tool set were transferred to the group of Wil van der Aalst at the Technical University of Eindhoven, The Netherlands [?]. New updates with improved functionality are made at a regular basis.

4. TIMED CPNS

5. SIDEBAR: STATE SPACES AND MODEL CHECKING

State space exploration is the main technique used for verifying behavioral properties of CPN models. In its basic form, a state space is a directed graph consisting of a node for each reachable marking (state) of the CPN model and edges corresponding to occurrences of enabled binding elements. From a constructed state space it is possible to automatically verify a wide range of standard behavioral properties of Petri nets (such as boundedness-, home-, and liveness properties). In addition, state spaces of CPN models can be used to perform model checking of behavioral properties expressed using temporal logics such as CTL and LTL. The main limitation of state space methods is the inherent state exploration problem which implies that state spaces in their

basic form are often too large to be with the available computing power. Many advanced techniques exist to combat the state explosion problem, and most of these can be applied also in the context of CPN models.

6. SIDEBAR: PERFORMANCE ANALYSIS

Simulation-based performance analysis is the main technique available for conducting quantitative analysis of timed CPN models. The basic idea of simulation-based performance analysis is to conduct a number of lengthy simulations of the CPN model under consideration. During the lengthy simulations data on, e.g., on queue length and delays, are collected. Data collection is based on the concept of monitors that observe simulations and write the extracted data into log files for post processing or directly compute key performance figures such as averages, standard deviation, and confidence intervals. In addition, batch simulations can be used to explore the parameter space of the model and conduct multiple simulation for each parameter in order to obtain statistically reliable results.

7. CPN TOOLS AND APPLICATIONS

The construction and analysis of CPN models have been supported by two generations of graphical computer tools. The first generation was the Design/CPN tool [?] which was developed starting in the mid 80's at Meta Software Corp. and later by the CPN Group at the Aarhus University. This was followed by CPN Tools [?] that has been developed since 2000 first by the CPN Group at Aarhus University and since 2009? by the XX group at the Technical University of Eindhoven. CPN Tools supports the editing and construction of CPN models, interactive and automatic simulation, state space-based model checking, and simulation-based performance analysis. Both Design/CPN and CPN Tools have been widely distributed tools and they have been applied for modelling and validation in a broad range of application domains. Below we provide some pointers to some selected applications within typical application domains. A more comprehensive list of example applications and domains can be found via [?].

Figure 7 provides a screenshot of CPN Tools. The user of CPN Tools works directly with the graphical representation of the CPN model. The graphical user interface (GUI) of CPN Tools has no conventional menu bars and pull-down menus, but is based on interaction techniques, such as *tool palettes* and *marking menus*. The rectangular area to the left is an *index*. It includes the *Tool box*, which is available for the user to manipulate the declarations and modules that constitute the CPN model. The *Tool box* includes tools for creating, copying, and cloning the basic elements of CP-nets. It also contains a wide selection of tools to manipulate the graphical layout and the appearance of the objects in the CPN model. The latter set of tools is very important in order to be able to create readable and graphically appealing CPN models. The remaining part of the screen is the *workspace*, which in this case contains four *binders* (the rectangular windows) and a circular pop-up menu.

Each binder holds a number of items which can be accessed by clicking the tabs at the top of the binder (only one item is visible at a time). There are two kinds of binders. One kind contains the elements of the CPN model, i.e., the mod-

ules and declarations. The other kind contains the tools which the user applies to construct and manipulate CPN models. The tools in a tool palette can be picked up with the mouse cursor and applied. In the example shown, one binder contains three modules named **Protocol**, **Sender**, and **Receiver**, while another binder contains a single module, named **Network**, together with the declaration of the colour set **NOxDATA**. The two remaining binders contain four different tool palettes to **Create** elements, change their **Style**, perform **Simulations**, and construct **State spaces**.

Items can be dragged from the index to the binders, and from one binder to another binder of the same kind. It is possible to position the same item in two different binders, for example, to view a module using two different zoom factors. A circular marking menu has been popped up on top of the bottom left binder. Marking menus are contextual menus that make it possible to select among the operations possible on a given object. In the case of Fig. ??, the marking menu gives the operations that can be performed on a port place object.

CPN Tools performs syntax and type checking, and error messages are provided to the user in a contextual manner next to the object causing the error. The syntax check and code generation are incremental and are performed in parallel with editing. This means that it is possible to execute parts of a CPN model even if the model is not complete, and that when parts of a CPN model are modified, a syntax check and code generation are performed only on the elements that depend on the parts that were modified. The main outcome of the code generation step is the *simulation code*. The simulation code contains the functions for inferring the set of enabled events in a given state of the CPN model, and for computing the state resulting from the occurrence (execution) of an enabled event in a given state.

CPN Tools supports two types of simulation: interactive and automatic. In an interactive simulation, the user is in complete control and determines the individual steps in the simulation, by selecting between the enabled events in the current state. CPN Tools shows the effect of executing a selected step in the graphical representation of the CPN model. In an automatic simulation the user specifies the number of steps that are to be executed and/or sets a number of stop criteria and breakpoints. The simulator then automatically executes the model without user interaction by making random choices between the enabled events in the states encountered. Only the resulting state is shown in the GUI. A *simulation report* can be saved, containing a specification of the steps that occurred during an automatic simulation. The simulator of CPN Tools exploits a number of advanced data structures for efficient simulation of large hierarchical CPN models. The simulator exploits the locality property of Petri nets to ensure that the number of steps executed per second in a simulation is independent of the size of the CPN model. This guarantees that simulation scales to large CPN models.

TALK/REFER TO THINGS ABOUT INDUSTRIAL APPLICATIONS IN THE END

Embedded Systems. Dalcotech or B and O

Process Scheduling. COAST

Internet Protocols. ERDP

Mobile Phone Software. NOKIA

Capacity Planning. HP

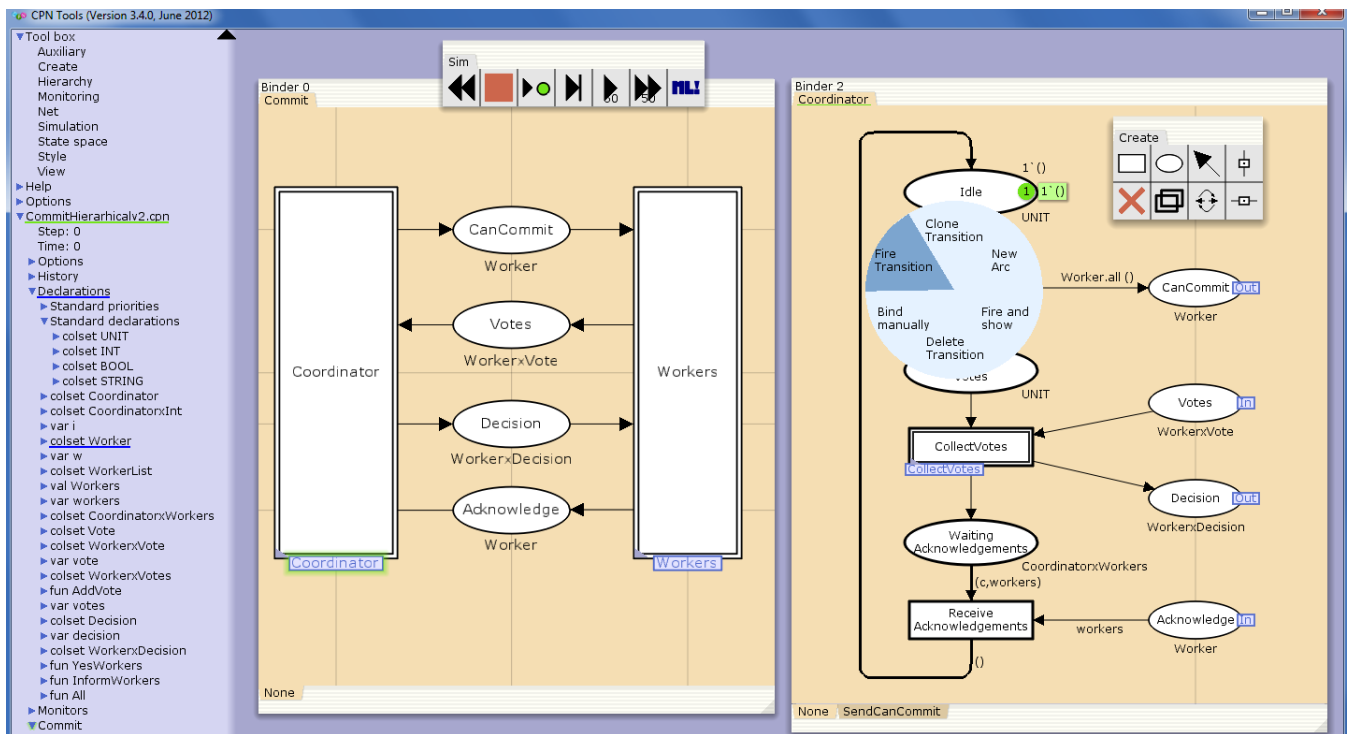


Figure 7: Two-phase Commit Protocol in CPN Tools